

Problem set 5

due Friday, December 8 at 11:59pm



Estimated time: 6 hours

Allocate about **1 hour per problem**, though some will take longer than others. You may need more time if programming is completely new to you, or less if you have some experience already.

Instructions Upload your .ipynb notebook to gradescope by 11:59pm on the due date.

- Note that each problem will be graded according to this [rubric](#). Solutions that include packages or functions not covered in this course will receive a score no higher than 2.
- You may collaborate with any of your classmates, but you must write your own code/solutions, understand all parts of the problem, and name your collaborators.
- You should also cite any outside sources you consulted, like Stack Overflow or ChatGPT, with a comment near the relevant lines of code (see example below). Recycled code that has not been cited will be considered plagiarism and receive a zero.

```
# code here was inspired by user2554330 on stack overflow:  
# https://stackoverflow.com/questions/69091812/is-everything-a-vector-in-r
```

As always, create a new colab R notebook. Include the title “Problem set 5”, your name, the date, and any collaborators somewhere at the top.

Dataset Information

You will work with the **verbs** dataset, part of the **languageR** package. This dataset is a simplified version of the **dative** dataset, which is described as follows:

Data describing the realization of the dative as NP or PP in the Switchboard corpus and the Treebank Wall Street Journal collection.

Dative alternation refers to the fact that some sentences can have two different structures that convey the same basic meaning. In the `verbs` dataset, we are looking at ditransitive verbs (such as ‘give’ or ‘offer’) which can realize the recipient in two ways: as a noun phrase (NP) “give you a book”, or as a prepositional phrase (PP) “give a book to you”. We want to investigate what factors in the data (`AnimacyOfRec`, `AnimacyOfTheme`, `LengthOfTheme`) impact whether the recipient is realized as a noun phrase (NP) or a prepositional phrase (PP) (`RealizationOfRec`).

Problem 1: Data Exploration

- Load the dataset and perform exploratory data analysis
- Examine the dataset’s structure (`glimpse`), summary statistics, and handle missing values if any.
- Visualize key relationships between variables using appropriate plots (scatter plots, histograms, etc.).

Problem 2: Model Specification

- Choose an appropriate type of model and explain why you’ve selected this type (e.g. regression, classification)
- Specify the model with an equation (e.g. $y = \beta_0 + \beta_1 x_1$). Note that you can create math equations in google colab text blocks by placing the equation between two dollar signs: `$y = \beta_0 + \beta_1 x_1$`
- Specify the model in R code (e.g. `y ~ x`)

Problem 3: Model Fitting

- Fit the model you specified in R
- Remember that the model fit also gives you predictions!
- What does the model tell us? How is realization of NP/PP impacted by the animacy of the theme (`AnimacyOfTheme`), the animacy of the recipient (`AnimacyOfRec`), and the length of the theme (`LengthOfTheme`)?
- Convert to probabilities where appropriate.

Problem 4: Model Accuracy

- Assess the accuracy of the model using cross-validation
- Take care to choose use `collect_metrics()` to report any accuracy metrics
- Add text explaining how accurate your model is and what this means.

Problem 5: Model Reliability

- Assess the reliability of your parameter estimates via bootstrapping
- Get a 95% confidence interval around the parameter estimates of your model.
- Visualize your bootstrapped distribution and shade the confidence interval.

Problem 6: Visualize your model

- Create a visualization to plot your model's fit (`geom_smooth()`)